

14. Critical Sections and Semaphores

SOA, Deemed to be University
ITER, Bhubanewar

Book(s)

Text Book(s)



Kay A. Robbins, & Steve Robbins

Unix™ Systems Programming

Communications, concurrency, and Treads

Pearson Education

Reference Book(s)



Brain W. Kernighan, & Rob Pike

The Unix Programming Environment

PHI

Introduction

- 👉 Programs that manage shared resources must execute portions of code called **critical sections** in a mutually exclusive manner.
- 👉 POSIX semaphores allow **processes** and threads to synchronize their actions.
- 👉 POSIX semaphores come in two forms: **named semaphores** and **unnamed semaphores**.

Semaphores

- ✎ A semaphore is an integer whose value is never allowed to fall below zero.
- ✎ Two operations can be performed on semaphores:
 - ✎ Increment the semaphore value by one (**sem_post**)
 - ✎ Decrement the semaphore value by one (**sem_wait**)
- ✎ If the value of a semaphore is currently zero, then a **sem_wait()** operation will block until the value becomes greater than zero.
- ✎ Other names for **wait** are **down**, **P** and **lock**.
- ✎ Other names for **signal** are **up**, **V**, **unlock** and **post**.

POSIX:SEM Unnamed Semaphores

- 👉 A POSIX:SEM semaphore is a variable of type `sem_t` with associated atomic operations for initializing, incrementing and decrementing its value.
- 👉 The difference between unnamed and named semaphores is analogous to the difference between ordinary pipes and named pipes (FIFOs).
- 👉 The following code segment declares a semaphore variable called `sem`.

```
#include <semaphore.h>

sem_t sem;
```

Unnamed Semaphore Initialization

```
#include <semaphore.h>

int sem_init(sem_t *sem, int pshared, unsigned
    value);
```

Returns:

- (1) If successful, `sem_init` initializes `sem`. `sem_init()` returns 0 on success.
- (2) If unsuccessful, `sem_init` returns -1 and sets `errno`.

The `sem_init` Parameters

- 👉: POSIX:SEM semaphores must be initialized before they are used.
- 👉: The `sem_init` function initializes the unnamed semaphore referenced by `sem` to `value`. The `value` parameter cannot be negative.

`pshared = 0`: means the semaphore can be used only by threads of the process that initializes the semaphore.

`pshared > 0`: If `pshared` is nonzero, any process that can access `sem` can use the semaphore.

`fork()`: simply forking a child after creating the semaphore does not provide access for the child. The child receives a copy of the semaphore, not the actual semaphore.

Unnamed Semaphore getvalue

```
#include <semaphore.h>

int sem_getvalue(sem_t *sem, int *sval);
```

Returns:

- (1) If successful, `sem_getvalue` returns 0.
- (2) If unsuccessful, `sem_getvalue` returns -1 and sets `errno`.

Unnamed Semaphore Wait Operation

```
#include <semaphore.h>

int sem_wait(sem_t *sem);

int sem_trywait(sem_t *sem);
```

Returns:

- (1) If successful, these functions **return** 0.
- (2) If unsuccessful, these functions **return** -1 and set `errno`.

Unnamed Semaphore Signal Operation

```
#include <semaphore.h>

int sem_post(sem_t *sem);
```

Returns:

- (1) If successful, `sem_post` returns 0.
- (2) If unsuccessful, `sem_post` returns -1 and sets `errno`.

Unnamed Semaphore Signal Operation

The `sem_destroy` function destroys a previously initialized unnamed semaphore referenced by the `sem` parameter.

```
#include <semaphore.h>

int sem_destroy(sem_t *sem);
```

Returns:

- (1) If successful, `sem_destroy` returns 0.
- (2) If unsuccessful, `sem_destroy` returns -1 and sets `errno`.

POSIX:SEM Named Semaphores

Introduction

- 👉 POSIX:SEM named semaphores can synchronize processes that do not share memory.
- 👉 Named semaphores have a name, a user ID, a group ID and permissions just as files do.
- 👉 A named semaphore is identified by a name of the form **/somename**. that is, a null-terminated string of characters (up to 251) consisting of an initial slash, followed by one or more characters, none of which are slashes.

Creating and Opening Named Semaphores

```
#include <fcntl.h>           /* For O_* constants */
#include <sys/stat.h>         /* For mode constants */
#include <semaphore.h>

sem_t *sem_open(const char *name, int oflag);

sem_t *sem_open(const char *name, int oflag, mode_t mode,
               unsigned int value);
```

Returns:

- (1) If successful, the `sem_open` function returns the address of the semaphore.
- (2) If unsuccessful, `sem_open` returns `SEM_FAILED` and sets `errno`.

Parameters of `sem_open`

- 👉 The **name** parameter is a string that identifies the semaphore by name.
- 👉 The **oflag** parameter determines whether the semaphore is created or just accessed by the function.
 - 👉 The **oflag** parameter is either 0, `O_CREAT`, or `O_CREAT | O_EXCL`.
 - 👉 If the `O_CREAT` is specified, the **sem_open** requires **two** more parameters: **a mode parameter of type `mode_t` giving the permissions** and **a value parameter of type `unsigned` giving the initial value of the semaphore**.
 - 👉 If both the `O_CREAT` and `O_EXCL` bits of **oflag** are set, the **sem_open** returns an `error` if the semaphore already exists.
 - 👉 If the semaphore already exists and `O_CREAT` is set but `O_EXCL` is not set, the semaphore ignores `O_CREAT` and the additional parameters.

Named Semaphore getvalue

```
#include <semaphore.h>

int sem_getvalue(sem_t *sem, int *sval);
```

Returns:

- (1) If successful, `sem_getvalue` returns 0.
- (2) If unsuccessful, `sem_getvalue` returns -1 and sets `errno`.

Named Semaphore Wait Operation

```
#include <semaphore.h>

int sem_wait(sem_t *sem);

int sem_trywait(sem_t *sem);
```

Returns:

- (1) If successful, these functions **return** 0.
- (2) If unsuccessful, these functions **return** -1 and set **errno**.

Named Semaphore Signal Operation

```
#include <semaphore.h>

int sem_post(sem_t *sem);
```

Returns:

- (1) If successful, `sem_post` returns 0.
- (2) If unsuccessful, `sem_post` returns -1 and sets `errno`.

Closing Named Semaphore

```
#include <semaphore.h>

int sem_close(sem_t *sem);
```

Returns:

- (1) If successful, `sem_close` returns 0.
- (2) If unsuccessful, `sem_close` returns -1 and sets `errno`

Unlinking Named Semaphore

```
#include <semaphore.h>

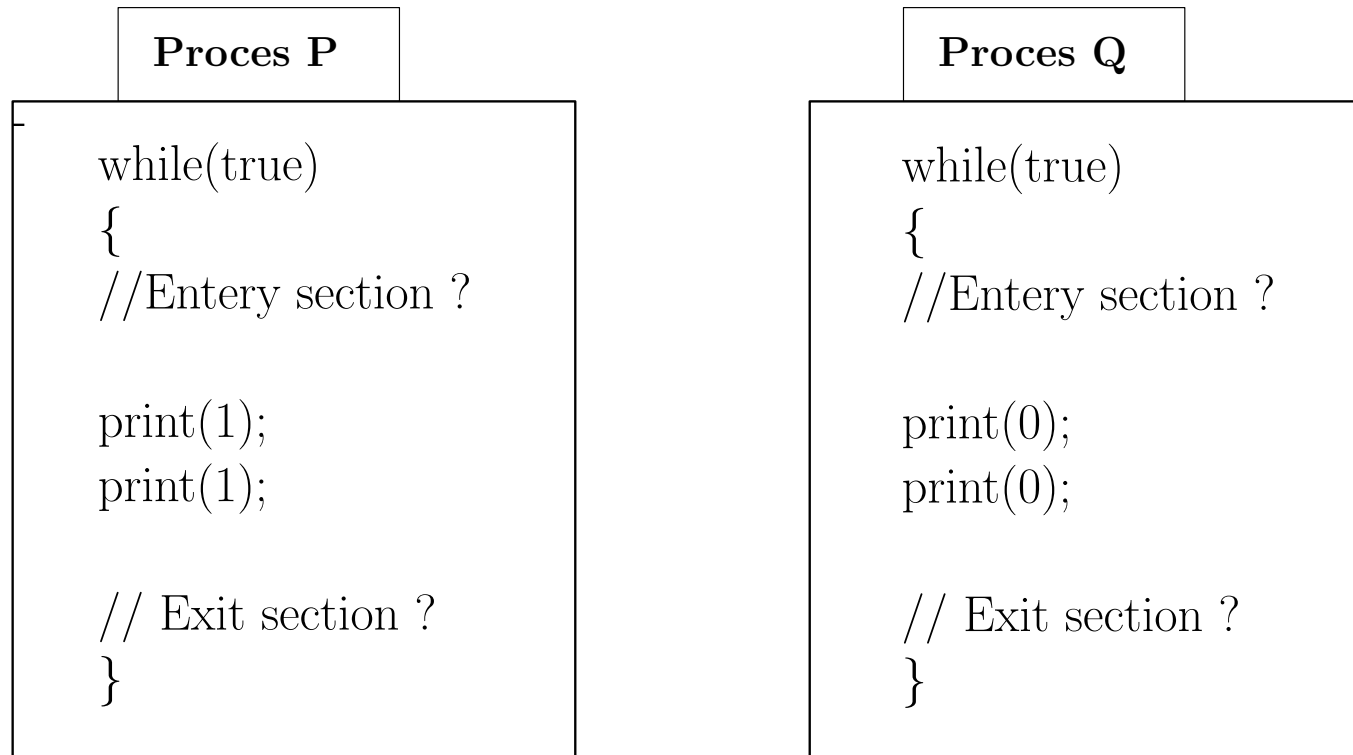
int sem_unlink(const char *name);
```

Returns:

- (1) If successful, `sem_unlink` returns 0.
- (2) If unsuccessful, `sem_unlink` returns -1 and sets `errno`.

DOS Mid-Sem Question, Feb-2019

Two concurrent processes **P** and **Q** are accessing their critical sections by using boolean variables **S** and **T** as follows;



Complete the entry section and exit section of process **P** and **Q** with suitable semaphore operations using the two boolean semaphores **S** and **T**. Also suggest the initial values of **S** and **T**, such that the execution of the processes will print the sequence 00110011.....