

UNIX IO

ITER, SOA (Deemed to be University), Bhubaneswar

By

Dr. TRILOK NATH PANDEY

DEPT. OF C.S.E



Kay A. Robbins, & Steve Robbins

Unix™ Systems Programming

Communications, concurrency, and Threads

Pearson Education

(Chapter 4)

Unix I/O:

- ☞ Unix uses a uniform device interface, through file descriptors, that allows the same I/O calls to be used for terminals, disks, tapes, audio and even network communication.
- ☞ There are basically five system calls used for I/O operations: **open, close, read, write and creat**

open System Call

- 👉 The `open` function associates a file descriptor with a file or physical device.

- 👉 SYNOPSIS:

```
#include <fcntl.h>
#include <sys/stat.h>
```

```
int open(const char *path, int oflag, ...);
```

(1) If successful, `open` returns a nonnegative integer representing the open file descriptor.

(2) If unsuccessful, `open` returns `-1` and sets `errno`.

- 👉 The `path` parameter of `open` points to the pathname of the file or device.
- 👉 The `oflag` parameter specifies status flags and access modes for the opened file.
- 👉 A third parameter must be included to specify access permissions if a file is created.

open () Examples

The following code segment opens a file my.dat for reading which belongs to the present working directory.

```
#include <fcntl.h>
#include <sys/stat.h>

int fd1;
fd1 = open("my.dat", O_RDONLY);
```

Modify the above code segment to open the file Users/jogs/ch4/file1.txt for reading

Perform OR the O_RDONLY and the O_CREAT flags.

```
myfd = open("Users/jogs/ch4/file1.txt", O_RDONLY
| O_CREAT);
```

O_CREAT is an additional flag which creates the file if the file does exist .

close () System Call

The close function has a single parameter, **filides**, representing the open file whose resources are to be released.

```
#include <unistd.h>
```

```
int close(int filides);
```

(1) If successful, close returns 0.

(2) If unsuccessful, close returns -1 and sets errno.

read System Call

- 👉 UNIX provides sequential access to files and other devices through the **read** and **write** functions.
- 👉 The **read** function attempts to retrieve **nbyte** bytes from the file or device represented by **filides** into the user variable **buf**.
- 👉 A large enough buffer must be provided to hold **nbyte** bytes of data.

👉 SYNOPSIS:

```
#include <unistd.h>
```

```
ssize_t read(int filides, void *buf, size_t nbyte);
```

- (1) If successful, **read** returns the number of bytes actually read.
- (2) If unsuccessful, **read** returns -1 and sets **errno**.

- 👉 The **ssize_t** data type is a signed integer data type used for the number of bytes read, or -1 if an error occurs.
- 👉 The **size_t** is an unsigned integer data type for the number of bytes to read.

Note: Read

- 👉 A **read** operation for a regular file may return fewer bytes than requested if, for example, it reached end-of-file before completely satisfying the request.
- 👉 A **read** operation for a regular file returns 0 to indicate end-of-file.
- 👉 When reading from a terminal, read returns 0 when the user enters an end-of-file character(CTRL+D).

write System Call






- 👉 The **write** function attempts to output **nbyte** bytes from the user buffer **buf** to the file represented by file descriptor **fildes**.
- 👉 SYNOPSIS:

```
#include <unistd.h>
```




```
ssize_t write(int fildes, const void *buf, size_t nbyte);
```

- (1) If successful, write returns the number of bytes actually written.
- (2) If unsuccessful, write returns -1 and sets errno.

File Descriptor

-  Files are designated with in C program either by **file pointers** or by **file descriptor**.
-  A file **descriptor** is an integer value that represents a file or device that is open.
-  It is an index into the process file descriptor table.
-  The file descriptor table is in the process user area and provides access to the system information for the associated file or device.
-  File pointers and file descriptors provide logical designations called *handles* for performing device-independent input and output.

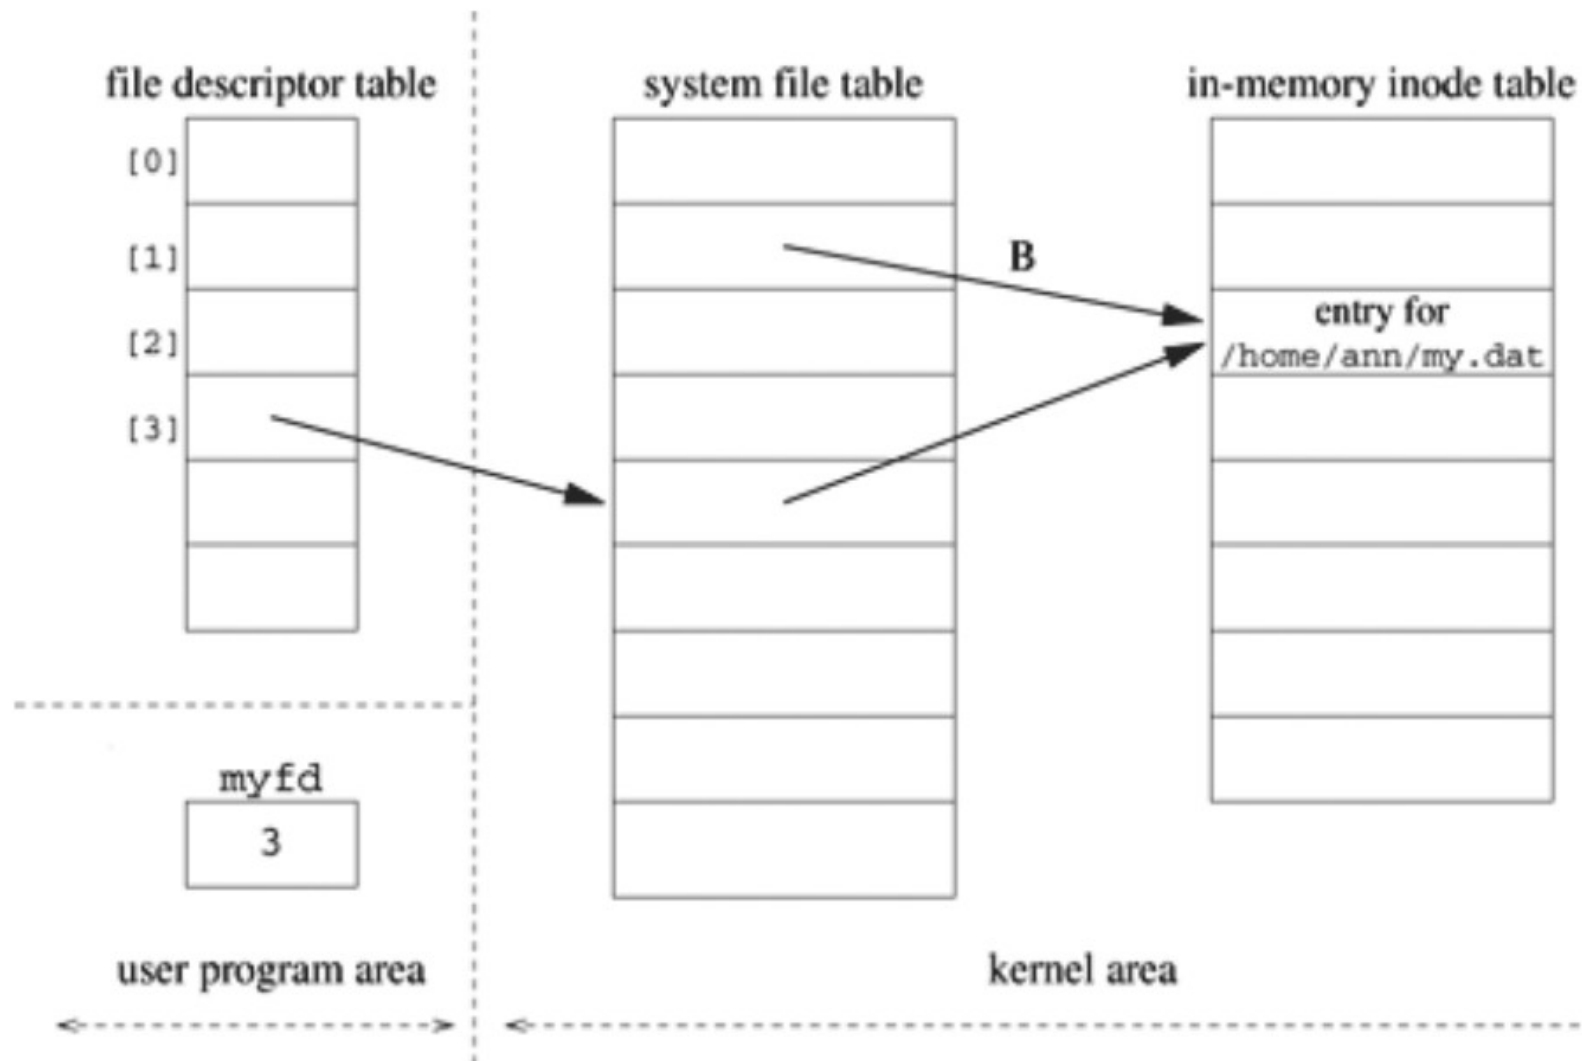
File Descriptor

-  The symbolic names for the **file pointers** that represent standard input, standard output and standard error are **`stdin`**, **`stdout`** and **`stderr`**, respectively. These symbolic names are defined in **`stdio.h`**.
-  The symbolic names for the **file descriptors** that represent standard input, standard output and standard error are **`STDIN_FILENO`**, **`STDOUT_FILENO`** and **`STDERR_FILENO`**, respectively. These symbolic names are defined in **`unistd.h`**.
-  The **numeric values** that represent standard input, standard output and standard error are 0, 1, and 2 respectively.

File Descriptor

A schematic of the file descriptor table after a program executes the following.

```
myfd = open("/home/ann/my.dat", O_RDONLY);
```



Examples of read() and write() system calls

The following program will read maximum 10 characters from standard input (e.g. keyboard) and write to standard output (e.g. display screen)

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/uio.h>
#include <unistd.h>
#include <fcntl.h>
int main()
{
    char buf[100];
    ssize_t bytesread;
    bytesread=read(STDIN_FILENO, buf, 10);
    write(STDOUT_FILENO, buf, bytesread);
    return 0;
}
```

Examples of read() and write() system calls cont...

Modify the previous program to read from a regular file (e.g. file1.txt) and write to standard output (e.g. display screen)

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/uio.h>
#include <unistd.h>
#include <fcntl.h>
int main()
{
    int fd1;
    char buf[10];
    ssize_t bytesread;
    fd1=open("file1.txt",O_RDONLY);
    bytesread=read(fd1, buf, 10);
    write(STDOUT_FILENO, buf, bytesread);
    close(fd1);
return 0;
}
```

Examples of read() and write() system calls Cont...

Write a program in C (copyfile.c) to copy the content of a regular file (file1.txt) to another regular file (file2.txt). Display the number of bytes read as well as number of bytes written.

```
#include<stdio.h>
#include<fcntl.h>
#include<unistd.h>
#define BLKSIZE 100
int main()
{
    int fd1,fd2;
    char buf[BLKSIZE];
    ssize_t bytesread,byteswrite;
    fd1=open("file1.txt",O_RDWR|O_CREAT, S_IRWXU);
    if(fd1==-1)
    {
        perror("Unable to open the file1.txt\n");
        return 0;
    }
    //printf("File Descriptor1=%d\n",fd1);
    bytesread=read(fd1,buf,BLKSIZE);
    close(fd1);
    fd2=open("file2.txt",O_WRONLY|O_CREAT, S_IRUSR|S_IWUSR);
    if(fd2==-1)
    {
        perror("Unable to open the file file2.txt\n");
        return 0;
    }
    //printf("File Descriptor2=%d\n",fd2);
    byteswrite=write(fd2,buf,bytesread);
    printf("bytesread=%d\nbyteswrite=%d\n",(int)bytesread,(int)byteswrite);
    return 0;
}
```

Examples of read() and write() system calls Cont...

Create a C file copyfile1.c by modifying the previous program and make an executable file (copy). Do the necessary steps to run it by name only like other unix commands:

Like \$ copy file1 file2 where file1 and file2 are source and destination files respectively.

```
#include<stdio.h>
#include<fcntl.h>
#include<unistd.h>
#define BLKSIZE 100
int main(int argc,char *argv[])
{
    int fd1,fd2;
    char buf[BLKSIZE];
    ssize_t bytesread,byteswrite;
    if(argc<3) {
        printf("Please provide source and destination files\n");
        return 0;
    }
    fd1=open(argv[1],O_RDWR|O_CREAT, S_IRWXU);
    if(fd1==-1) {
        perror("Unable to open the source file\n");
        return 0;
    }
    bytesread=read(fd1,buf,BLKSIZE);
    close(fd1);
    fd2=open(argv[2],O_WRONLY|O_CREAT, S_IRUSR|S_IWUSR);
    if(fd2==-1)
    {
        perror("Unable to open the file destination file\n");
        return 0;
    }
    byteswrite=write(fd2,buf,bytesread);
    if( byteswrite<=0)
        printf("0 bytes copied\n");
    return 0;
}
```


Examples of read() and write() system calls Cont...

Step 1: make executable file (copy) of the program copyfile.c

```
$ gcc copyfile1.c -o copy
```

Step 2: Include the absolute path of the file “copy” into the values of environment variable PATH

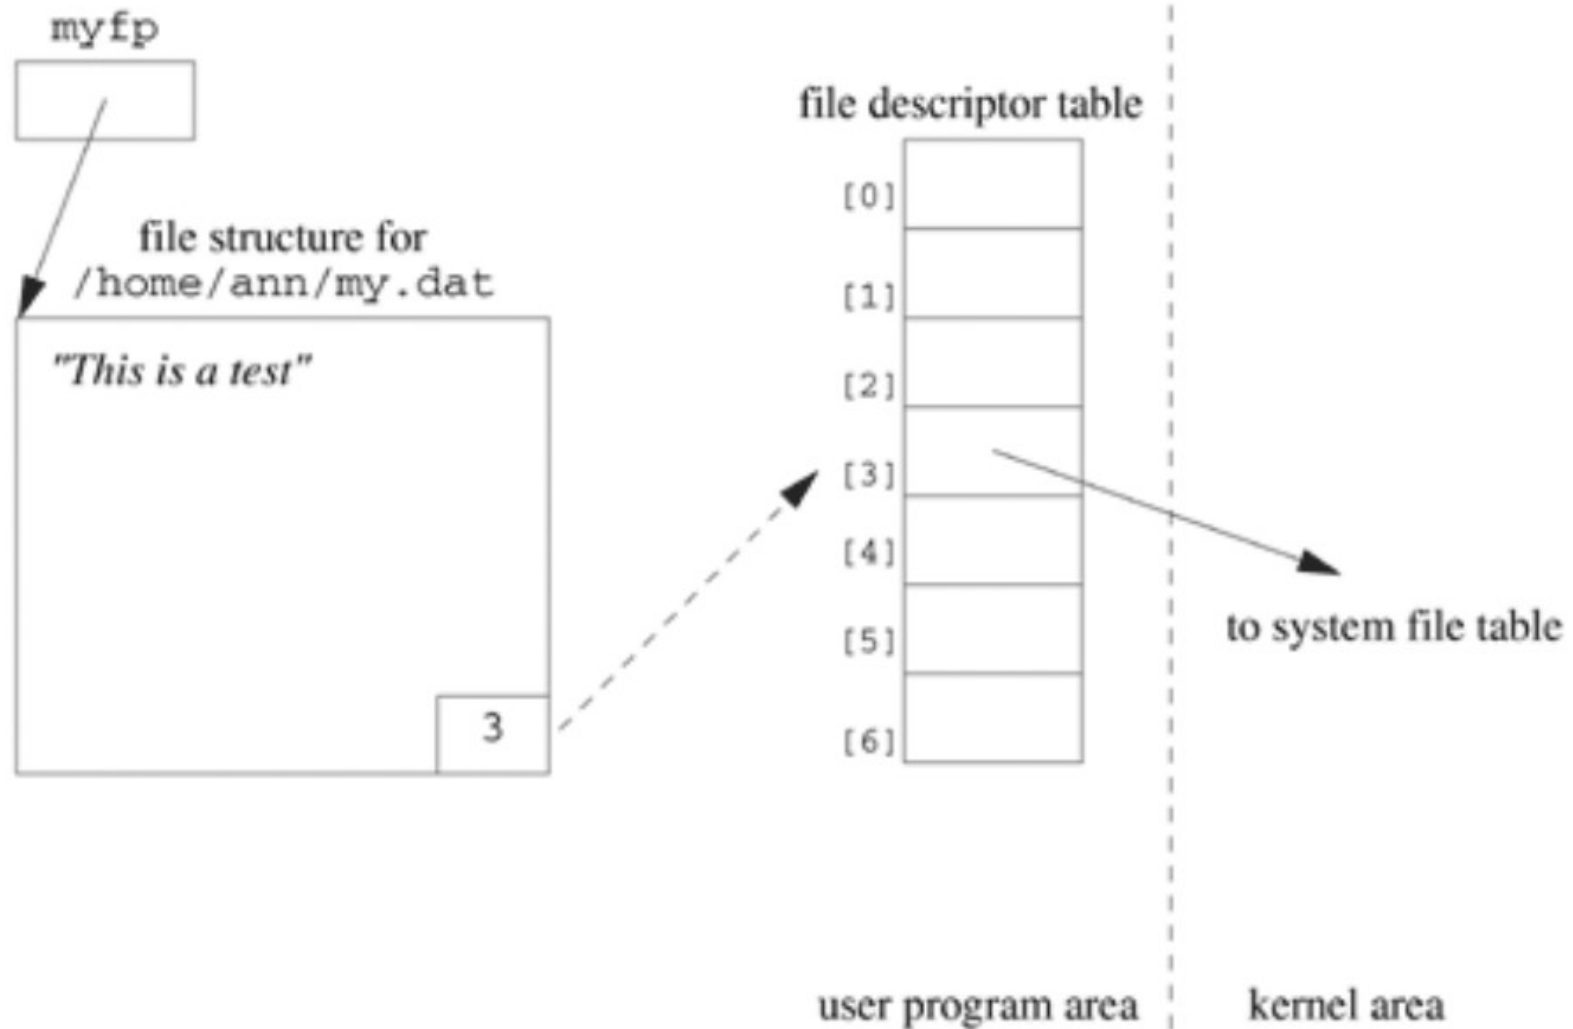
```
export PATH=$PATH:/Users/jogs/ch4/Lec3
```

Step 3: run the executable file ‘copy’

```
$ copy file1.txt file2.txt
```

File Pointer

```
FILE *myfp;  
if ((myfp = fopen("/home/ann/my.dat", "w")) == NULL)  
    perror("Failed to open /home/ann/my.dat");  
else  
    fprintf(myfp, "This is a test");
```



File descriptor's value from file pointer

```
#include<stdio.h>
int main()
{
    FILE *myfp;
    int fd;
    myfp=fopen("Trial.txt","w");
    if(myfp==NULL)
    {
        perror("Opening Error");
        return 1;
    }
    /* To get the file descriptor value */
    fd=fileno(myfp); /* fileno() is a library function */
    printf("File descriptor=%d\n",fd);
    return 0;
}
```

oflag Argument Setting

 The POSIX values for the **access mode flags** in **oflag** argument:

1. **O_RDONLY** : read-only access
2. **O_WRONLY** : write-only access
3. **O_RDWR** : read-write-only access

Note: specify exactly one of above designating read-only, write-only or read-write access.

 The **oflag** argument is also constructed by taking the bitwise OR (|) of the desired combination of the access mode and the **additional flags**.

 The additional flags:

1. **O_APPEND**
2. **O_CREAT**
3. **O_EXCL**
4. **O_NOCTTY**
5. **O_NONBLOCK**
6. **O_TRUNC**

Additional Flags

O_APPEND: The O_APPEND flag causes the file offset to be moved to the end of the file before a write, allowing you to add to an existing file.

O_CREAT: The O_CREAT flag causes a file to be created if it doesn't already exist. If O_CREAT flag is included, a third argument to `open()` must be passed to designate the permissions.

O_EXCL: If you want to avoid writing over an existing file, use the combination O_CREAT | O_EXCL. This combination returns an error if the file already exists.

O_NOCTTY: The O_NOCTTY flag prevents an opened device from becoming a controlling terminal.


O_NONBLOCK: The O_NONBLOCK flag controls whether the `open()` returns immediately or blocks until the device is ready.

O_TRUNC: O_TRUNC truncates the length of a regular file opened for writing to 0.

Third argument to `open()`

- 👉 Each file has three classes associated with it: a user (or owner), a group and everybody else (others).
- 👉 The possible permissions or privileges are read(r), write(w) and execute(x). These privileges are specified separately for the user, the group and others.
- 👉 When you open a file with the `O_CREAT` flag, you must specify the permissions as the third argument to `open` in a mask of type `mode_t`.
- 👉 POSIX defines symbolic names for masks corresponding to the permission bits. These names are defined in **`sys/stat.h`**

POSIX symbolic names for file permissions








Symbol	meaning
S_IRUSR S_IWUSR R S_IXUSR S_IRWXU	read by owner write by owner execute by owner read, write, execute by owner
S_IRGRP S_IWGRP S_IXGRP S_IRWXG	read by group write by group execute by group read, write, execute by group
S_IROTH S_IWOTH H S_IXOTH S_IRWXO	read by others write by others execute by others read, write, execute by others
S_ISUID S_ISGID	set user ID on execution set group ID on execution

Filters

- ✎ A filter reads from standard input, performs a transformation, and outputs the result to standard output.
- ✎ Filters write their error messages to standard error.
- ✎ **Examples:** `head`, `tail`, `more`, `sort`, `grep`, `sed`, and `awk` etc.
- ✎ **`cat` as a filter:** The `cat` command takes a list of filenames as command-line arguments, reads each of the files in succession, and echoes the contents of each file to standard output. However, if no input file is specified, `cat` takes its input from standard input and writes its results to standard output. In this case, `cat` behaves like a filter.

Redirection

-  a file descriptor is an index into the file descriptor table of that process.
-  Each entry in the file descriptor table points to an entry in the system file table, which is created when the file is opened.
-  A program can modify the file descriptor table entry so that it points to a different entry in the system file table. This action is known as **redirection**.
-  Most shells interpret the greater than character (>) on the command line as redirection of standard output and the less than character (<) as redirection of standard input.
-  **Example:** `$ cat > myfile.txt`, redirects standard output to `myfile.txt` with `>`.

dup () & dup2 () : Duplicate a File Descriptor

```
#include <unistd.h>

int dup(int fildes);
```

Return:

- (1) dup() creates a copy of the file descriptor fildes.
- (2) dup() uses the lowest-numbered unused descriptor **for** the new descriptor returned
- (3) -1 on error

```
#include <unistd.h>

fd=open("read.c",O_RDONLY);
nfd=dup(fd);
printf("Duplicate fd=%d\n",nfd);
```

/ duplicates the file descriptor fd */*

nfd is the lowest-numbered unused file descriptor. It is the duplicate of **fd**.

Example: dup ()

Duplicate the standard output file descriptor to write onto a file descriptor `fd`

```
/*duplicating STDOUT_FILENO to a file descriptor fd */
int main()
{
    int fd;  fd=open("duptest.txt",O_WRONLY|O_CREAT|
        O_TRUNC,S_IRUSR|
        S_IWUSR|S_IRGRP);
    printf("File descriptor: fd=%d\n",fd);

    dup(STDOUT_FILENO);/* save descriptor STDOUT_FILENO */
    close(1);           /* closing 1, creates an empty slot */
    dup(fd);
    close(fd);          /* duplicate fd to standard output */
    write(STDOUT_FILENO,"USP\n",4);
    write(STDOUT_FILENO,"DOS\n",4);
    return 0;
}
```

Run the code, then open the file: `$ cat duptest.txt`. Data is now written onto the file instead of monitor

dup2 () : Duplicate a File Descriptor

```
#include <unistd.h>

int dup2(int fildes, int fildes2);
```

Return:

- (1) On success, dup2 returns the file descriptor value that was duplicated.
- (2) -1 on error

Example: dup2 ()

Duplicate the standard output file descriptor to write onto a file descriptor `fd`

```
/*duplicating STDOUT_FILENO to a file descriptor fd */

#define CREATE_FLAGS (O_WRONLY | O_CREAT | O_APPEND)
#define CREATE_MODE (S_IRUSR | S_IWUSR | S_IRGRP | S_IROTH)
int main(void) {
    int fd;
    fd = open("dup2test.txt", CREATE_FLAGS, CREATE_MODE);
    if (dup2(fd, STDOUT_FILENO) == -1) {
        perror("Failed to redirect standard output");
        return 1;
    }
    close(fd);
    write(STDOUT_FILENO, "OK", 2);
    return 0;
}
```

Run the code, then open the file: `$ cat dup2test.txt`. Data is now written onto the file instead of monitor

dup () vs dup2 ()

```
dup2 (fd1 , fd2) ;
```

Is equivalent to

```
close (fd2) ;  
dup (fd1) ;
```

Thank You